

CSE301 – Computer Organization

Lecture 5 – Instruction-Level Parallelism (ILP)

Instruction-Level Parallelism (ILP) refers to the ability of a processor to **execute multiple instructions concurrently** by exploiting parallelism within a single program. The idea is to overlap or run independent instructions at the same time to increase performance.

There are **two main types** of ILP:

1. Static ILP (Compiler-Controlled ILP)

In static ILP, instruction parallelism is extracted **at compile time**. The compiler analyzes the program, detects independent instructions, and schedules them to execute in parallel.

Characteristics:

- The compiler determines which instructions can issue together in the same cycle.
- Parallelism is fixed before runtime.
- Very dependent on compiler quality.

Static ILP Architectures

1. **SISD (Single Instruction, Single Data)** Traditional sequential computers.
2. **SIMD (Single Instruction, Multiple Data)** One instruction operates on many data elements simultaneously. Common in GPUs and vector processors.
3. **MIMD (Multiple Instruction, Multiple Data)** Multiple processors execute different instructions on different data streams.
4. **MISD (Multiple Instruction, Single Data)** Rare; instructions operate in parallel on the same data.

VLIW (Very Long Instruction Word)

A key form of static ILP.

- The compiler packs multiple independent operations into a single long instruction.
- Each part of the instruction is executed by a different functional unit in the same clock cycle.
- No dynamic scheduling hardware → simpler CPU but complex compiler.

Note: In static ILP, the compiler is responsible for ensuring that issued instructions have **no data dependencies**.

2. Dynamic ILP (Hardware-Controlled ILP)

In dynamic ILP, the processor extracts parallelism **during runtime**. The CPU itself decides:

- How many instructions to fetch per cycle
- How many to issue and execute in parallel
- How many results to write back (commit) per cycle

This leads to **superscalar processors**, which fetch and issue **multiple instructions every cycle**.

Dynamic ILP depends heavily on **Tomasulo's Algorithm**, used for:

- Out-of-order execution
- Detecting dependencies at runtime

- Avoiding stalls
- Register renaming

Superscalar Execution

A superscalar processor uses wide instruction pipelines that allow:

- Fetching multiple instructions per cycle
- Dispatching them to different functional units
- Executing them out of order
- Committing results in order

The number of instructions fetched/issued per cycle defines the CPU's **issue width**.

Tomasulo's Algorithm (Dynamic Scheduling)

Tomasulo's algorithm enables **out-of-order execution** while preserving correct program behavior.

Key ideas:

1. **Instructions share a common fetch/decode/commit pipeline**, done in-order.
2. **Instructions are issued to functional units out-of-order**, based on:
 - The type of operation
 - Availability of operands
 - Availability of the functional unit
3. **Reservation stations** hold instructions until they are ready to execute.
4. **Common Data Bus (CDB)** broadcasts results so waiting instructions can receive operands immediately without reading the register file.

This eliminates many RAW hazards and significantly increases ILP.




Figure 1: Superscaler processor

Register Renaming

A critical part of Tomasulo's algorithm.

Why register renaming?

To eliminate **false dependencies**:

- **WAR (Write After Read)**
- **WAW (Write After Write)**

These dependencies occur because multiple instructions use the same architectural registers.

How renaming works

- Each architectural register (e.g., `$t0`) is mapped to a **physical register**.
- New instructions get **new physical registers** for their destination operands.
- This prevents later instructions from overwriting a register before earlier instructions read it.

Example

Without renaming:

```
add r1, r2, r3    ; writes to r1
sub r1, r4, r5    ; also writes to r1 → WAW
```

With renaming:

add → P7

sub → P9

Even though both target "r1", they map to different *physical* registers → no conflict.

This allows:

- Out-of-order execution
- Parallel execution of independent instructions
- Elimination of false hazards (WAR, WAW)

Here is a polished, clean, and lecture-ready section you can add under the **Instruction-Level Parallelism (ILP)** topic.

Loop Unrolling

Loop Unrolling is a compiler optimization technique used to **increase instruction-level parallelism (ILP)** by reducing loop control overhead and exposing more independent operations to the processor.

The main idea is to **duplicate the loop body multiple times** and reduce the number of iterations. This creates more straight-line (non-branch) code, allowing the processor—especially superscalar or pipelined CPUs—to execute more instructions in parallel.

Why Loop Unrolling Improves Performance

1. **Reduces branch penalty** Fewer branch instructions → fewer control hazards.
 2. **Increases ILP** More instructions become visible to:
 - the compiler (in static ILP)
 - or the hardware scheduler (in dynamic ILP)
 3. **More opportunities for:**
 - instruction scheduling
 - register renaming
 - forwarding and pipelining
 - hiding latency (e.g., load latencies)
-

Basic Example

Original loop:

```
for (i = 0; i < 4; i++) {  
    a[i] = a[i] + 1;  
}
```

Unrolled by a factor of 4:

```
a[0] = a[0] + 1;  
a[1] = a[1] + 1;
```

```
a[2] = a[2] + 1;  
a[3] = a[3] + 1;
```

Benefits:

- Branch executes once instead of four times
- Independent instructions → can issue in parallel
- Better utilization of functional units

MIPS Example

Original:

```
Loop:  
    lw    $t0, 0($s1)  
    addi $t0, $t0, 1  
    sw    $t0, 0($s1)  
    addi $s1, $s1, 4  
    addi $s2, $s2, -1  
    bne   $s2, $zero, Loop
```

Unrolled by 2:

```
Loop:  
    lw    $t0, 0($s1)  
    addi $t0, $t0, 1  
    sw    $t0, 0($s1)  
  
    lw    $t1, 4($s1)  
    addi $t1, $t1, 1  
    sw    $t1, 4($s1)  
  
    addi $s1, $s1, 8  
    addi $s2, $s2, -2  
    bne   $s2, $zero, Loop
```

Now the branch runs only half as often, and the CPU can attempt to overlap the two independent iterations.